

# Cassandra Assignment 2 - SWEN 432

Zoltan Debre - 300360191

Original repository and progress history: <https://github.com/zoltan-nz/cassandra-exercise>

## Question 1

(2 marks)

- Use `ccm` to make a single data center Cassandra cluster having 5 nodes call it `single_dc`
- Start the cluster
- Run the `ccm node1 ring` command
- Save the output of the ring command for future use and show it in the answer to the question.

```
$ ccm create single_dc --nodes=5
$ ccm start
```

```
$ ccm status -v
```

```
Cluster: 'single_dc'
```

```
-----
node1: UP
```

```
  auto_bootstrap=False
  thrift=('127.0.0.1', 9160)
  binary=('127.0.0.1', 9042)
  storage=('127.0.0.1', 7000)
  jmx_port=7100
  remote_debug_port=0
  byteman_port=0
  initial_token=-9223372036854775808
  pid=41519
```

```
node3: UP
```

```
  auto_bootstrap=False
  thrift=('127.0.0.3', 9160)
  binary=('127.0.0.3', 9042)
  storage=('127.0.0.3', 7000)
  jmx_port=7300
  remote_debug_port=0
  byteman_port=0
  initial_token=-1844674407370955162
  pid=41547
```

```
node2: UP
```

```
  auto_bootstrap=False
  thrift=('127.0.0.2', 9160)
  binary=('127.0.0.2', 9042)
  storage=('127.0.0.2', 7000)
  jmx_port=7200
  remote_debug_port=0
  byteman_port=0
  initial_token=-553402322112865485
  pid=41588
```

```
node5: UP
```

```
  auto_bootstrap=False
  thrift=('127.0.0.5', 9160)
  binary=('127.0.0.5', 9042)
  storage=('127.0.0.5', 7000)
  jmx_port=7500
  remote_debug_port=0
  byteman_port=0
  initial_token=553402322112865484
  pid=41629
```

```
node4: UP
  auto_bootstrap=False
  thrift=('127.0.0.4', 9160)
  binary=('127.0.0.4', 9042)
  storage=('127.0.0.4', 7000)
  jmx_port=7400
  remote_debug_port=0
  byteman_port=0
  initial_token=1844674407370955161
  pid=41665
```

```
$ ccm node1 ring
```

```
Datacenter: datacenter1
```

```
=====
```

Address	Rack	Status	State	Load	Owns	Token
127.0.0.1	rack1	Up	Normal	98.97 KiB	40.00%	5534023222112865484
127.0.0.2	rack1	Up	Normal	98.98 KiB	40.00%	-9223372036854775808
127.0.0.3	rack1	Up	Normal	98.98 KiB	40.00%	-5534023222112865485
127.0.0.4	rack1	Up	Normal	98.98 KiB	40.00%	-1844674407370955162
127.0.0.5	rack1	Up	Normal	98.98 KiB	40.00%	1844674407370955161
						5534023222112865484

## Question 2

(14 marks)

- a) (2 marks) What is the setting of the `endpoint_snitch` property?

Location of `node1`'s `cassandra.yaml`: `~/ccm/single_dc/node1/conf/cassandra.yaml`

(You can find a copy of `cassandra.yaml` in `node1` sub-folder.)

```
endpoint_snitch: SimpleSnitch
```

- b) (6 Marks) What is the value of the `initial_token` property?

```
initial_token: -9223372036854775808
```

- Which Cassandra component has calculated it?

default partitioner:

```
Murmur3Partitioner
```

- Is there any relationship between `initial_token` property value and the output of the `ccm node1 ring` command?

Yes.

The first node from the `ccm node1 ring` output:

127.0.0.1	rack1	Up	Normal	98.97 KiB	40.00%	-9223372036854775808
-----------	-------	----	--------	-----------	--------	----------------------

The hash value of the first node is the same as the `initial_token`.

- c) (2 marks) What is the setting of the `partitioner` property?

```
partitioner: org.apache.cassandra.dht.Murmur3Partitioner
```

- d) (4 marks) What is the setting of the `rpc_address` property?

```
rpc_address: 127.0.0.1
```

- Is there any relationship between `rpc_address` property value and the output of the `ccm node1 ring` command?

Yes, this value is the ip address of the first node

```
127.0.0.1 rack1 Up Normal 98.97 KiB 40.00% -9223372036854775808
```

### Question 3

(2 marks) Consider the `cassandra.topology.properties` file of `node1` and comment on the relationship between file's content and the output of the `ccm node1 ring` command.

Location of the file: `~/ccm/single_dc/node1/conf/cassandra-topology.properties`  
(Copy saved in `node1` sub-folder.)

The main content of the property file:

```
# Cassandra Node IP=Data Center:Rack
192.168.1.100=DC1:RAC1
192.168.2.200=DC2:RAC2

10.0.0.10=DC1:RAC1
10.0.0.11=DC1:RAC1
10.0.0.12=DC1:RAC2

10.20.114.10=DC2:RAC1
10.20.114.11=DC2:RAC1

10.21.119.13=DC3:RAC1
10.21.119.10=DC3:RAC1

10.0.0.13=DC1:RAC2
10.21.119.14=DC3:RAC2
10.20.114.15=DC2:RAC2

# default for unknown nodes
default=DC1:r1
```

Because of using a single data center (Simple Snitch) we don't see any relation between this file and our `ccm node1 ring` output. Simple Snitch does not recognize data center or rack information. Our property file content is just dummy data, examples, not related to our node.

### Question 4

(8 marks)

a) (3marks) Connect to `cqlsh` prompt and create a keyspace with the name `ass2`. Replication strategy should be `simple`, and the replication factor equal `3`. In your answer, show your keyspace declaration.

```
$ ccm node1 cqlsh
cqlsh> CREATE KEYSPACE IF NOT EXISTS ass2 with replication = { 'class' : 'SimpleStrategy', 'replication_factor' :
```

b) (5marks)

The following files:

```
table_declarations.cql
data_point_data.txt
driver_data.txt
time_table_data.txt
vehicle_data.txt
```

are given on the course Assignments page. The file `table_declarations.cql` contains create table statements, while the other files contain comma separated table data. Use these files, and `SOURCE` and `COPY` `cqlsh` commands to implement a version of the train time table data base. In your answer show the results of running the `cqlsh` command `describe tables` and of running `select` statements on each table for a row of your choice.

```
$ ccm node1 cqlsh
cqlsh> SOURCE './table_declarations.cql'
cqlsh> DESCRIBE TABLES;
```

#### Keyspace system\_schema

```
-----
tables      triggers    views    keyspaces  dropped_columns
functions   aggregates  indexes  types      columns
```

#### Keyspace system\_auth

```
-----
resource_role_permissions_index  role_permissions  role_members  roles
```

#### Keyspace system

```
-----
available_ranges    peers          batchlog        transferred_ranges
batches             compaction_history  size_estimates  hints
prepared_statements sstable_activity  built_views
"IndexInfo"         peer_events      range_xfers
views_builds_in_progress  paxos           local
```

#### Keyspace system\_distributed

```
-----
repair_history  view_build_status  parent_repair_history
```

#### Keyspace system\_traces

```
-----
events  sessions
```

#### Keyspace ass2

```
-----
time_table  data_point  driver  vehicle
```

```
cqlsh> COPY ass2.time_table (line_name, service_no, time, distance, latitude, longitude, stop) FROM './time_table'
Using 7 child processes
```

```
Starting copy of ass2.time_table with columns [line_name, service_no, time, distance, latitude, longitude, stop].
Processed: 30 rows; Rate:      44 rows/s; Avg. rate:      66 rows/s
30 rows imported from 1 files in 0.458 seconds (0 skipped).
cqlsh> SELECT * FROM ass2.time_table;
```

line_name	service_no	time	distance	latitude	longitude	stop
Melling	3	807	13.7	-41.2036	174.9054	Melling
Melling	3	801	11.4	-41.2118	174.89	Western Hutt
Melling	3	754	8.3	-41.227	174.8851	Petone
Melling	3	741	0	-41.2865	174.7762	Wellington
Hutt Valley Line	1	650	34.3	-41.1244	175.0708	Upper Hutt
Hutt Valley Line	1	642	26.5	-41.1479	175.0122	Silverstream
Hutt Valley Line	1	634	19	-41.1798	174.9608	Taita
Hutt Valley Line	1	629	15.8	-41.2024	174.9423	Naenae
Hutt Valley Line	1	625	13.3	-41.2092	174.9081	Waterloo
Hutt Valley Line	1	622	11	-41.2204	174.9081	Woburn
Hutt Valley Line	1	617	8.3	-41.227	174.8851	Petone
Hutt Valley Line	1	605	0	-41.2865	174.7762	Wellington
Waikanae	5	1139	62.8	-40.8755	175.0668	Waikanae
Waikanae	5	1118	51.3	-40.9142	175.0084	Paraparaumu
Waikanae	5	1059	33.1	-40.9881	174.951	Paekakariki
Waikanae	5	1042	15.9	-41.1339	174.8406	Porirua
Waikanae	5	1025	0	-41.2865	174.7762	Wellington
Hutt Valley Line	11	2025	34.3	-41.1244	175.0708	Upper Hutt
Hutt Valley Line	11	2019	26.5	-41.1479	175.0122	Silverstream
Hutt Valley Line	11	2010	19	-41.1798	174.9608	Taita
Hutt Valley Line	11	2001	15.8	-41.2024	174.9423	Naenae
Hutt Valley Line	11	1955	13.3	-41.2092	174.9081	Waterloo
Hutt Valley Line	11	1952	11	-41.2204	174.9081	Woburn
Hutt Valley Line	11	1947	8.3	-41.227	174.8851	Petone
Hutt Valley Line	11	1935	0	-41.2865	174.7762	Wellington
Hutt Valley Line	2	1045	34.3	-41.2865	174.7762	Wellington
Hutt Valley Line	2	1033	26	-41.227	174.8851	Petone
Hutt Valley Line	2	1025	21	-41.2092	174.9081	Waterloo
Hutt Valley Line	2	1015	15.3	-41.1798	174.9608	Taita
Hutt Valley Line	2	1000	0	-41.1244	175.0708	Upper Hutt

(30 rows)

```
cqlsh> COPY ass2.data_point FROM './data_point_data.txt';
Using 7 child processes
```

```
Starting copy of ass2.data_point with columns [line_name, service_no, date, sequence, latitude, longitude, speed].
Processed: 5 rows; Rate:      8 rows/s; Avg. rate:      11 rows/s
5 rows imported from 1 files in 0.437 seconds (0 skipped).
cqlsh> SELECT * FROM ass2.data_point;
```

line_name	service_no	date	sequence	latitude	longitude	speed
Hutt Valey Line	2	20160326	2016-03-25 21:07:40.000000+0000	-41.2012	175	70.1
Hutt Valey Line	2	20160326	2016-03-25 21:02:10.000000+0000	-41.1255	175.07	40.5
Hutt Valey Line	2	20160326	2016-03-24 21:27:10.000000+0000	-41.2262	174.77	29.1
Hutt Valey Line	2	20150322	2015-03-21 21:44:10.000000+0000	-41.2862	174.7759	9.1
Hutt Valey Line	2	20160322	2016-03-21 21:37:50.000000+0000	-41.2272	174.77	29.1

(5 rows)

```
cqlsh> COPY ass2.driver FROM './driver_data.txt';
Using 7 child processes
```

```
Starting copy of ass2.driver with columns [driver_name, current_position, email, mobile, password, skill].
Processed: 6 rows; Rate:      10 rows/s; Avg. rate:      14 rows/s
6 rows imported from 1 files in 0.430 seconds (0 skipped).
cqlsh> SELECT * FROM ass2.driver;
```

driver_name	current_position	email	mobile	password	skill
fred	Taita	fred@ecs.vuw.ac.nz	2799797	f00f	{'Ganz Mavag', 'Guliver
jane	Waikanae	jane@ecs.vuw.ac.nz	2131131	jj77	{'Matangi
ann	not available	ann@ecs.vuw.ac.nz	21998877	aaaa	{'Matangi
milan	Upper Hutt	milan@ecs.vuw.ac.nz	211111	mm77	{'Matangi
pondy	Wellington	pondy@ecs.vuw.ac.nz	214455	pd66	{'Guliver', 'Matangi
pavle	Upper Hutt	pmogin@ecs.vuw.ac.nz	213344	pm33	{'Ganz Mavag', 'Guliver', 'Matangi

(6 rows)

```
cqlsh> COPY ass2.vehicle FROM './vehicle_data.txt';
Using 7 child processes
```

```
Starting copy of ass2.vehicle with columns [vehicle_id, status, type].
Processed: 6 rows; Rate:      10 rows/s; Avg. rate:      14 rows/s
6 rows imported from 1 files in 0.433 seconds (0 skipped).
cqlsh> SELECT * FROM ass2.vehicle;
```

vehicle_id	status	type
KW3300	Wellington	Matangi
FP3003	out of order	Guliver
FA3456	in_use	Matangi
FP8899	Upper Hutt	Matangi
FA4864	maintenance	Matangi
FA1122	Upper Hutt	Ganz Mavag

(6 rows)

## Question 5

(10 marks) To answer this question, you will need to use the `getendpoints nodetool` command.

a) (1 mark) Find the nodes storing data of driver `pavle`. In your answer, show the output of the `getendpoints nodetool` command. Let us call these nodes `node_a`, `node_b`, and `node_c`.

```
$ ccm node1 nodetool getendpoints ass2 driver pavle
```

```
127.0.0.1
127.0.0.2
127.0.0.3
```

b) (3 marks)

Connect to `cqlsh` prompt using a node that is not in the set `{node_a, node_b, node_c}`.

```
$ ccm node4 cqlsh
```

Set the consistency level to ALL and read data of the driver pavle.

```
cqlsh> CONSISTENCY;
Current consistency level is ONE.
cqlsh> CONSISTENCY ALL;
Consistency level set to ALL.
cqlsh> SELECT * FROM ass2.driver WHERE driver_name='pavle';

driver_name | current_position | email                | mobile | password | skill
-----+-----+-----+-----+-----+-----
pavle       | Upper Hutt      | pmogin@ecs.vuw.ac.nz | 213344 | pm33     | {'Ganz Mavag', 'Guliver', 'Matangi'}

(1 rows)
```

Stop `node_a`, connect to `cqlsh`, set the consistency level to ALL and read pavle's data again. What have you learned?

```
$ ccm node1 stop
$ ccm node4 cqlsh
cqlsh> CONSISTENCY ALL;
Consistency level set to ALL.
cqlsh> SELECT * FROM ass2.driver WHERE driver_name='pavle';
NoHostAvailable:
```

`CONSISTENCY ALL` in Read Consistency Levels means that Cassandra returns the record after all replicas have responded. The read operation will fail if a replica does not respond. Exactly this happened in our case.

c) (3 marks)

With `node_a` still being stopped, set the consistency level to QUORUM and read pavle's data.

```
cqlsh> CONSISTENCY QUORUM ;
Consistency level set to QUORUM.
cqlsh> SELECT * FROM ass2.driver WHERE driver_name='pavle';

driver_name | current_position | email                | mobile | password | skill
-----+-----+-----+-----+-----+-----
pavle       | Upper Hutt      | pmogin@ecs.vuw.ac.nz | 213344 | pm33     | {'Ganz Mavag', 'Guliver', 'Matangi'}

(1 rows)
```

Stop `node_b`, connect to `cqlsh`, set the consistency level to QUORUM and read pavle's data again. What have you learned

```
$ ccm node2 stop
$ ccm node4 cqlsh
cqlsh> CONSISTENCY QUORUM;
Consistency level set to QUORUM.
cqlsh> SELECT * FROM ass2.driver WHERE driver_name='pavle';
NoHostAvailable:
```

The `CONSISTENCY QUORUM` means that Cassandra will return the record after a quorum of replicas from all datacenters has responded. We have one datacenter at this stage, the replication factor is 3, so at least 2 should respond. We got results when only one node was dead, but no responses when 2 were dead from 3 nodes.

d) (3 marks)

With `node_a` and `node_b` still being stopped, set the consistency level to ONE and read pavle's data.

```
cqlsh> CONSISTENCY ONE ;
Consistency level set to ONE.
cqlsh> SELECT * FROM ass2.driver WHERE driver_name='pavle';
```

driver_name	current_position	email	mobile	password	skill
pavle	Upper Hutt	pmogin@ecs.vuw.ac.nz	213344	pm33	{'Ganz Mavag', 'Guliver', 'Matangi'}

(1 rows)

Stop node\_c, connect to cqlsh, and read pavle's data again. What have you learned

```
$ ccm node3 stop
$ ccm node4 cqlsh
cqlsh> CONSISTENCY ONE;
Consistency level set to ONE.
cqlsh> SELECT * FROM ass2.driver WHERE driver_name='pavle';
NoHostAvailable:
```

**CONSISTENCY ONE** : Returns a response from the closest replica, as determined by the snitch. We have one data center and 3 replicas. In the first case we still had one node available, so one replica was still existed. But after stopping node\_c does not left any live replica, so no response.

## Question 6

(15 marks)

You are asked to find those nodes of the `single_dc` Cassandra cluster that store replicas of driver `eileen`. Very soon you realized that all `ccm` commands and `nodetool` commands, including `ccm start`, `ccm stop`, `ccm status`, `ccm nodei cqlsh` and so on, work properly except the command `ccm nodei nodetool getendpoints ass2 driver eileen`. Despite that, you have devised a procedure to find the nodes requested. In your answer, describe the procedure and show how you have applied it.

First of all, we don't have driver with `eileen` in our database. We can check it with the following query.

```
$ ccm start
$ ccm switch single_dc
$ ccm status

Cluster: 'single_dc'
-----
node1: UP
node3: UP
node2: UP
node5: UP
node4: UP

$ ccm node1 cqlsh -e "use ass2; select * from driver where driver_name = 'eileen';"

 driver_name | current_position | email | mobile | password | skill
-----+-----+-----+-----+-----+-----
(0 rows)

$ ccm node1 cqlsh -e "INSERT INTO ass2.driver (driver_name, current_position, email, mobile, password, skill) VALUES ('eileen', 'Wellington', 'eileen@ecs.vuw.ac.nz', '555444', 'abcd123', {'Guliver', 'Matangi'});"
[applied]
-----
True

ccm node1 cqlsh -e "use ass2; select * from driver where driver_name = 'eileen';"

 driver_name | current_position | email | mobile | password | skill
-----+-----+-----+-----+-----+-----
eileen | Wellington | eileen@ecs.vuw.ac.nz | 555444 | abcd123 | {'Guliver', 'Matangi'}

(1 rows)
```

Without using `nodetool`, we can find nodes which stores our record, if we close all other nodes except one. Using **CONSISTENCY ONE**, we get back our record if that node stores our requested data.

The following bash script can help us to iterate through on all nodes and run the query. (`q6-node-finder.sh`)

```
nodes=('node1' 'node2' 'node3' 'node4' 'node5');
for node in "${nodes[@]}; do
    ccm switch single_dc
    ccm stop
    ccm ${node} start
    echo "Active node: ${node}"
    ccm status
    ccm ${node} cqlsh -e "use ass2; consistency one; select * from driver where driver_name='eileen';"
done
```

Or we can run manually this one line for each node:

```
$ ccm switch single_dc; ccm stop; ccm node1 start; ccm status; ccm node1 cqlsh -e "use ass2; consistency one; select * from driver where driver_name='eileen';"
```

Cluster: 'single\_dc'

```
-----
node1: UP
node3: DOWN
node2: DOWN
node5: DOWN
node4: DOWN
```

driver_name	current_position	email	mobile	password	skill
eileen	Wellington	eileen@ecs.vuw.ac.nz	555444	abcd123	{'Guliver', 'Matangi'}

(1 rows)

```
$ ccm switch single_dc; ccm stop; ccm node2 start; ccm status; ccm node2 cqlsh -e "use ass2; consistency one; select * from driver where driver_name='eileen';"
```

Cluster: 'single\_dc'

```
-----
node1: DOWN
node3: DOWN
node2: UP
node5: DOWN
node4: DOWN
```

Consistency level set to ONE.

driver_name	current_position	email	mobile	password	skill
eileen	Wellington	eileen@ecs.vuw.ac.nz	555444	abcd123	{'Guliver', 'Matangi'}

(1 rows)

```
$ ccm switch single_dc; ccm stop; ccm node3 start; ccm status; ccm node3 cqlsh -e "use ass2; consistency one; select * from driver where driver_name='eileen';"
```

Cluster: 'single\_dc'

```
-----
node1: DOWN
node3: UP
node2: DOWN
node5: DOWN
node4: DOWN
```

Consistency level set to ONE.

NoHostAvailable:

```
$ ccm switch single_dc; ccm stop; ccm node4 start; ccm status; ccm node4 cqlsh -e "use ass2; consistency one; select * from driver where driver_name='eileen';"
```

Cluster: 'single\_dc'

```
-----
node1: DOWN
node3: DOWN
node2: DOWN
node5: DOWN
node4: UP
```

Consistency level set to ONE.

NoHostAvailable:

```
$ ccm switch single_dc; ccm stop; ccm node5 start; ccm status; ccm node5 cqlsh -e "use ass2; consistency one; select * from driver where driver_name='eileen';"
```



```
Cluster: 'single_dc'
```

```
-----  
node1: DOWN  
node3: DOWN  
node2: DOWN  
node5: UP  
node4: DOWN
```

Consistency level set to ONE.

```
driver_name | current_position | email | mobile | password | skill  
-----  
eileen | Wellington | eileen@ecs.vuw.ac.nz | 555444 | abcd123 | {'Guliver', 'Matangi'}  
  
(1 rows)
```

We can see, that our record exists on `node1`, `node2` and `node5`.

We can test this with our restricted command:

```
$ ccm start; ccm node1 nodetool getendpoints ass2 driver eileen  
127.0.0.5  
127.0.0.1  
127.0.0.2
```

Update:

I found a more elegant solution when I solved Question 15. Using the token function.

```
$ ccm start  
$ ccm node1 cqlsh -e "USE ass2; SELECT driver_name, TOKEN(driver_name) FROM driver WHERE driver_name='eileen'"  
  
driver_name | system.token(driver_name)  
-----  
eileen | 2694043365177161046  
  
(1 rows)
```

Comparing our above token number with previously printed ring token thresholds, we can see, that `eileen`'s token is above `node4`'s threshold but smaller than `node5`. It means, that the primary node is `node5`. In the ring the following nodes, which will store replicas, are `node1` and `node2`. Again, the answer is `node1`, `node2` and `node5`.

```
Datacenter: datacenter1  
=====
```

Address	Rack	Status	State	Load	Owns	Token
...						
127.0.0.4	rack1	Up	Normal	98.98 KiB	40.00%	1844674407370955161
127.0.0.5	rack1	Up	Normal	98.98 KiB	40.00%	5534023222112865484

## Question 7

(15 marks)

Assume the following situation:

The data of the driver `james` should be stored on `node4`, `node5`, and `node1`.

A client (say `c0`) connected to `node3` and sent a request to write `james`'s data.

In the moment of running the statement `insert into driver (driver_name, password) values ('james', '7007');` `node4` was down.

Writing succeeded.

In the next moment `node5` and `node1` went down and the `node4` started.

A client (say c1) connected to `cqlsh` prompt via `node3` and sent the following read statement: `select driver_name, password from driver where driver_name = 'james';`

The read result was:

driver_name	password
james	7007

Repeat the experiment described above. Name and briefly explain Cassandra mechanism that made succeeding of the select statement above possible.

We can see where cassandra would like to store our record. Because the primary key in `driver` table is the `driver_name`, we can list our nodes. And it is really the `node1`, `node4` and `node5`.

```
$ ccm start
$ ccm node1 nodetool getendpoints ass2 driver james

127.0.0.4
127.0.0.5
127.0.0.1

$ ccm status

Cluster: 'single_dc'
-----
node1: UP
node3: UP
node2: UP
node5: UP
node4: UP
```

Simulate `node4` is down.

```
$ ccm node4 stop
$ ccm status
Cluster: 'single_dc'
-----
node1: UP
node3: UP
node2: UP
node5: UP
node4: DOWN
```

We may insert our new row with the following command.

```
$ ccm node3 cqlsh -e "use ass2; insert into driver (driver_name, password) values ('james', '7007');"
```

However, if we don't change our consistency level than our experiment will fail. The default consistency level `ONE`. In this case Cassandra will not replicate our record to `node4` when it starts. We have to switch at least to `CONSISTENCY QUORUM`.

We have to use the following command to insert our data.

```
$ ccm node3 cqlsh -e "use ass2; consistency quorum; insert into driver (driver_name, password) values ('james', '7007');"
```

Simulate `node1`, `node5` are down and `node4` is back.

```
$ ccm node1 stop; ccm node5 stop; ccm node4 start
$ ccm status
Cluster: 'single_dc'
-----
node1: DOWN
node3: UP
node2: UP
node5: DOWN
node4: UP
```

```
$ ccm node3 cqlsh -e "use ass2; select driver_name, password from driver where driver_name='james';"

driver_name | password
-----+-----
james | 7007

(1 rows)
```

When we run our insert command with quorum consistency, Cassandra write our record at least in two nodes, plus it will write in log and in the memtable. When our `node1` and `node5` is stoped and `node4` came back, Cassandra inserted our record from memtable in `node4` .

Cassandra uses gossip process to track states of nodes. It helps to determine which node is up or down and when it can replicate a missing data.

## Question 8

Lost...

## Question 9

(3 marks)

Use ccm to make a Cassandra cluster spanning two datacenters. The cluster name should be `multi_dc` . Cassandra will automatically assign default names `dc1` and `dc2` to datacenters. The cluster `multi_dc` uses 5 nodes in `dc1` and 4 nodes in `dc2` . Start the cluster and run the `ccm ring` command. Save the output of the ring command for future use and show it in the answer to the question.

```
$ ccm create -n 5:4 -s multi_dc
$ ccm switch multi_dc
$ ccm start
$ ccm status

Cluster: 'multi_dc'
-----
node9: UP
node8: UP
node1: UP
node3: UP
node2: UP
node5: UP
node4: UP
node7: UP
node6: UP

$ ccm node1 ring

Datacenter: dc1
=====
Address      Rack      Status State   Load           Owns           Token
127.0.0.1    r1        Up      Normal  98.97 KiB      25.00%        5534023222112865484
127.0.0.2    r1        Up      Normal  98.96 KiB      20.00%        -9223372036854775808
127.0.0.3    r1        Up      Normal  98.97 KiB      20.00%        -5534023222112865485
127.0.0.4    r1        Up      Normal  98.97 KiB      20.00%        -1844674407370955162
127.0.0.5    r1        Up      Normal  98.96 KiB      20.00%        1844674407370955161
127.0.0.5    r1        Up      Normal  98.96 KiB      20.00%        5534023222112865484

Datacenter: dc2
=====
Address      Rack      Status State   Load           Owns           Token
127.0.0.6    r1        Up      Normal  98.96 KiB      20.00%        4611686018427388004
127.0.0.7    r1        Up      Normal  98.97 KiB      25.00%        -9223372036854775708
127.0.0.8    r1        Up      Normal  98.94 KiB      25.00%        -4611686018427387804
127.0.0.8    r1        Up      Normal  98.94 KiB      25.00%        100
127.0.0.9    r1        Up      Normal  98.97 KiB      25.00%        4611686018427388004
```

## Question 10

(4 marks)

Consider the `cassandra.yaml` file of `node1`. What is the setting of the `endpoint_snitch` property? If you find it different to the setting in the case of the `single_dc` cluster, explain briefly why it is different.

(Please find the `cassandra.yaml` in the `multi_dc` folder.)

```
endpoint_snitch: org.apache.cassandra.locator.PropertyFileSnitch
```

In case of `single_dc` our `endpoint_snitch` value was `SimpleSnitch`. In case of `multi_dc` Cassandra uses `PropertyFileSnitch` protocol.

The `SimpleSnitch` is used only for single-datacenter deployments. `PropertyFileSnitch` determines the location of nodes by rack and datacenter. It uses the network details located in the `cassandra-topology.properties` file (copied in `multi_dc` folder).

## Question 11

(4 marks)

Consider the `cassandra.topology.properties` file of `node1` and comment on the relationship between file's content and the output of the `ccm node1 ring` command.

Content of `cassandra.topology.properties` :

```
default=dc1:r1
127.0.0.1=dc1:r1
127.0.0.2=dc1:r1
127.0.0.3=dc1:r1
127.0.0.4=dc1:r1
127.0.0.5=dc1:r1
127.0.0.6=dc2:r1
127.0.0.7=dc2:r1
127.0.0.8=dc2:r1
127.0.0.9=dc2:r1
```

`cassandra.topology.properties` file contains the same mapping as we can list with `ccm node1 ring`. We see in both cases that we have two data centers and which IP address, which node belongs to `dc1` or to `dc2` datacenter.

Both list the rack numbers also. In our case, we use only one-one rack.

## Question 12

(2 marks)

Create a keyspace with the name `ass2` having network topology replication strategy and a replication factor of 3 for both `dc1` and `dc2` datacenters. In your answer, show your keyspace declaration.

```
cqlsh> CREATE KEYSPACE IF NOT EXISTS ass2 WITH replication = {'class': 'NetworkTopologyStrategy', 'dc1': 3, 'dc2': 3}
```

## Question 13

(3 marks)

Use `SOURCE` and `COPY` `cqlsh` commands and the following files:

```
table_declarations.cql
driver_data.txt
time_table_data.txt
```

to implement a version of the train time table data base. You need to populate only `driver` and `time_table` tables by data. In your answer show the results of running the `cqlsh` command `describe tables` and of running CQL `select` statements

on `driver` and `time_table` for a row of your choice.

```
$ ccm node1 cqlsh -e "use ass2; SOURCE './table_declarations.cql';"
```

```
$ ccm node1 cqlsh
cqlsh> DESCRIBE tables;
```

Keyspace system\_schema

```
-----
tables      triggers    views    keyspaces  dropped_columns
functions   aggregates  indexes  types      columns
```

Keyspace system\_auth

```
-----
resource_role_permissions_index  role_permissions  role_members  roles
```

Keyspace system

```
-----
available_ranges      peers      batchlog      transferred_ranges
batches               compaction_history  size_estimates  hints
prepared_statements   sstable_activity  built_views
"IndexInfo"           peer_events      range_xfers
views_builds_in_progress  paxos          local
```

Keyspace system\_distributed

```
-----
repair_history  view_build_status  parent_repair_history
```

Keyspace system\_traces

```
-----
events  sessions
```

Keyspace ass2

```
-----
time_table  data_point  driver  vehicle
```

```
cqlsh> USE ass2;
```

```
cqlsh:ass2> COPY driver FROM './driver_data.txt';
```

Using 7 child processes

Starting copy of ass2.driver with columns [driver\_name, current\_position, email, mobile, password, skill].

Processed: 6 rows; Rate: 6 rows/s; Avg. rate: 9 rows/s

6 rows imported from 1 files in 0.664 seconds (0 skipped).

```
cqlsh:ass2> SELECT * FROM driver;
```

driver_name	current_position	email	mobile	password	skill
fred	Taita	fred@ecs.vuw.ac.nz	2799797	f00f	{'Ganz Mavag', 'Guliver
jane	Waikanae	jane@ecs.vuw.ac.nz	2131131	jj77	{'Matangi
ann	not available	ann@ecs.vuw.ac.nz	21998877	aaaa	{'Matangi
milan	Upper Hutt	milan@ecs.vuw.ac.nz	211111	mm77	{'Matangi
pondy	Wellington	pondy@ecs.vuw.ac.nz	214455	pd66	{'Guliver', 'Matangi
pavle	Upper Hutt	pmogin@ecs.vuw.ac.nz	213344	pm33	{'Ganz Mavag', 'Guliver', 'Matangi

(6 rows)

```
cqlsh:ass2> COPY time_table FROM './time_table_data.txt';
```

Using 7 child processes

Starting copy of ass2.time\_table with columns [line\_name, service\_no, time, distance, latitude, longitude, stop].

Processed: 30 rows; Rate: 47 rows/s; Avg. rate: 86 rows/s

30 rows imported from 1 files in 0.348 seconds (0 skipped).

```
cqlsh:ass2> SELECT * FROM time_table;
```

line_name	service_no	time	distance	latitude	longitude	stop
Melling	3	807	13.7	-41.2036	174.9054	Melling
Melling	3	801	11.4	-41.2118	174.89	Western Hutt
Melling	3	754	8.3	-41.227	174.8851	Petone
Melling	3	741	0	-41.2865	174.7762	Wellington
Hutt Valley Line	1	650	34.3	-41.1244	175.0708	Upper Hutt
Hutt Valley Line	1	642	26.5	-41.1479	175.0122	Silverstream
Hutt Valley Line	1	634	19	-41.1798	174.9608	Taita
Hutt Valley Line	1	629	15.8	-41.2024	174.9423	Naenae
Hutt Valley Line	1	625	13.3	-41.2092	174.9081	Waterloo

Hutt Valley Line	1	622	11	-41.2204	174.9081	Woburn
Hutt Valley Line	1	617	8.3	-41.227	174.8851	Petone
Hutt Valley Line	1	605	0	-41.2865	174.7762	Wellington
Waikanae	5	1139	62.8	-40.8755	175.0668	Waikanae
Waikanae	5	1118	51.3	-40.9142	175.0084	Paraparaumu
Waikanae	5	1059	33.1	-40.9881	174.951	Paekakariki
Waikanae	5	1042	15.9	-41.1339	174.8406	Porirua
Waikanae	5	1025	0	-41.2865	174.7762	Wellington
Hutt Valley Line	11	2025	34.3	-41.1244	175.0708	Upper Hutt
Hutt Valley Line	11	2019	26.5	-41.1479	175.0122	Silverstream
Hutt Valley Line	11	2010	19	-41.1798	174.9608	Taita
Hutt Valley Line	11	2001	15.8	-41.2024	174.9423	Naenae
Hutt Valley Line	11	1955	13.3	-41.2092	174.9081	Waterloo
Hutt Valley Line	11	1952	11	-41.2204	174.9081	Woburn
Hutt Valley Line	11	1947	8.3	-41.227	174.8851	Petone
Hutt Valley Line	11	1935	0	-41.2865	174.7762	Wellington
Hutt Valley Line	2	1045	34.3	-41.2865	174.7762	Wellington
Hutt Valley Line	2	1033	26	-41.227	174.8851	Petone
Hutt Valley Line	2	1025	21	-41.2092	174.9081	Waterloo
Hutt Valley Line	2	1015	15.3	-41.1798	174.9608	Taita
Hutt Valley Line	2	1000	0	-41.1244	175.0708	Upper Hutt

(30 rows)

## Question 14

(8 marks)

Find nodes storing data of the driver `pavle`. Let these nodes be `node_a`, `node_b`, `node_c`, `node_d`, `node_e`, and `node_f`, where  $a < b < c < d < e < f$ .

```
$ ccm node1 nodetool getendpoints ass2 driver pavle
```

```
127.0.0.1
127.0.0.2
127.0.0.3
127.0.0.6
127.0.0.7
127.0.0.8
```

l. (4 marks)

Connect to ass2 keyspace.

```
$ ccm node1 cqlsh
Connected to multi_dc at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.10 | CQL spec 3.4.4 | Native protocol v4]
Use HELP for help.
cqlsh> USE ass2;
cqlsh:ass2>
```

Run the statement: `select driver_name, password from driver where driver_name = 'pavle';` under consistency levels:

- `quorum`

```
cqlsh:ass2> CONSISTENCY QUORUM;
Consistency level set to QUORUM.
cqlsh:ass2> SELECT driver_name, password FROM driver WHERE driver_name = 'pavle';
```

```
driver_name | password
-----+-----
pavle      | pm33
```

(1 rows)

- `each_quorum`

```
cqlsh:ass2> CONSISTENCY EACH_QUORUM;
Consistency level set to EACH_QUORUM.
cqlsh:ass2> SELECT driver_name, password FROM driver WHERE driver_name = 'pavle';

driver_name | password
-----+-----
pavle | pm33

(1 rows)
```

Run the select statement under consistency level `local_quorum` once for `dc1` being local, and once for `dc2` being local.

- `local_quorum` (on node1)

```
$ ccm node1 cqlsh -e "USE ass2; CONSISTENCY local_quorum; SELECT driver_name, password FROM driver WHERE driver_name = 'pavle';"
Consistency level set to LOCAL_QUORUM.

driver_name | password
-----+-----
pavle | pm33

(1 rows)
```

```
$ ccm node6 cqlsh -e "USE ass2; CONSISTENCY local_quorum; SELECT driver_name, password FROM driver WHERE driver_name = 'pavle';"
Consistency level set to LOCAL_QUORUM.

driver_name | password
-----+-----
pavle | pm33

(1 rows)
```

Our all nodes are live. We have at least 2 nodes from 3 in each datacenter. Quorum means that Cassandra returns the record after a quorum of replicas has responded from any data center. Based on the documentation `each_quorum` is not supported for reads, however our query worked and it will be clear effect when we stop nodes in the next task. In case of `local_quorum` Cassandra returns the record after a quorum of replicas in the current data center. We can avoid latency of inter-data center communication.

II. (4 marks)

Use `ccm` to stop `node_e` and `node_f`. Connect to `ass2` keyspace.

```
$ ccm node7 stop
$ ccm node8 stop
$ ccm status

Cluster: 'multi_dc'
-----
node9: UP
node8: DOWN
node1: UP
node3: UP
node2: UP
node5: UP
node4: UP
node7: DOWN
node6: UP
```

Run the statement `select driver_name, password from driver where driver_name = 'pavle';` under consistency levels: `quorum`, `each_quorum`, and `local_quorum`. Run the select statement under consistency level `local_quorum` once for `dc1` being local, and once for `dc2` being local. In your answer to the question, show results of your experiments and describe briefly what you have learned.

```
$ ccm node1 cqlsh -e "USE ass2; CONSISTENCY quorum; SELECT driver_name, password FROM driver WHERE driver_name = 'pavle';"
Consistency level set to QUORUM.

driver_name | password
-----+-----
```

```
pavle | pm33

(1 rows)

$ ccm node6 cqlsh -e "USE ass2; CONSISTENCY quorum; SELECT driver_name, password FROM driver WHERE driver_name =
Consistency level set to QUORUM.

driver_name | password
-----+-----
pavle | pm33

(1 rows)
```

We got proper respond, because at least 2 nodes live on one of the cluster. Doesn't matter to which cluster the client connects.

```
$ ccm node1 cqlsh -e "USE ass2; CONSISTENCY each_quorum; SELECT driver_name, password FROM driver WHERE driver_name =
Consistency level set to EACH_QUORUM.
<stdin>:1:NoHostAvailable:

$ ccm node6 cqlsh -e "USE ass2; CONSISTENCY each_quorum; SELECT driver_name, password FROM driver WHERE driver_name =
Consistency level set to EACH_QUORUM.
<stdin>:1:NoHostAvailable:
```

Our request will fail, because `each_quorum` consistency expect that at least 2 nodes are active in each data center. It forces strong consistency. We should use this level in multiple datacenter clusters to strictly maintain consistency at the same level in each datacenter and if we want a read to fail when a datacenter is down and the QUORUM cannot be reached on that datacenter. Exactly this happened with us in this case.

```
$ ccm node1 cqlsh -e "USE ass2; CONSISTENCY local_quorum; SELECT driver_name, password FROM driver WHERE driver_name =
Consistency level set to LOCAL_QUORUM.

driver_name | password
-----+-----
pavle | pm33

(1 rows)

$ ccm node6 cqlsh -e "USE ass2; CONSISTENCY local_quorum; SELECT driver_name, password FROM driver WHERE driver_name =
Consistency level set to LOCAL_QUORUM.
<stdin>:1:NoHostAvailable:
```

We can see, that `local_quorum` expects that we have at least 2 active nodes with replica from the connected datacenter. In the first case, when we connected to `node1` we had enough nodes active in `dc1`. In the second case, when we connected `node6` which is in `dc2`, there was not more live replica, so the quorum is not satisfied.

## Question 15

(10 marks)

You are asked to find those nodes of the `multi_dc` Cassandra cluster that store replicas of the `time_table` table row

line_name	service_no	time	distance	latitude	longitude	stop
Hutt Valley Line	2	1045	34.3	-41.2865	174.7762	Wellington

Very soon you realized that all `ccm` and `nodetool` commands, except `ccm nodei cqlsh`, do not work. So, you are unable to use: `ccm stop`, `ccm status`, `ccm start`, `ccm nodei ring` and so on, including the command `ccm nodei nodetool getendpoints ass2 time_table <key>`.

Despite that, you have devised a procedure to find the nodes requested. In your answer, describe the procedure and show how you have applied it.

Hint: Luckily, you have saved the output of the `ccm nodei ring` command and `cqlsh` prompt is still working.



We can use the primary key for finding the default node. Using the token function with primary keys of `time_table` we can list the token number of each row. The token number determines the default node of a record. Because our replication level is 3, the default node and the following 2 nodes will store our record.

```
$ ccm node1 cqlsh -e "USE ass2; SELECT line_name, service_no, time, token(line_name, service_no) as t FROM time_table;"

line_name      | service_no | time | t
-----+-----+-----+-----
Melling        | 3          | 807  | -7474942320664480980
Melling        | 3          | 801  | -7474942320664480980
Melling        | 3          | 754  | -7474942320664480980
Melling        | 3          | 741  | -7474942320664480980
Hutt Valley Line | 1          | 650  | -6012480106752428297
Hutt Valley Line | 1          | 642  | -6012480106752428297
Hutt Valley Line | 1          | 634  | -6012480106752428297
Hutt Valley Line | 1          | 629  | -6012480106752428297
Hutt Valley Line | 1          | 625  | -6012480106752428297
Hutt Valley Line | 1          | 622  | -6012480106752428297
Hutt Valley Line | 1          | 617  | -6012480106752428297
Hutt Valley Line | 1          | 605  | -6012480106752428297
Waikanae       | 5          | 1139 | -5905794062536720418
Waikanae       | 5          | 1118 | -5905794062536720418
Waikanae       | 5          | 1059 | -5905794062536720418
Waikanae       | 5          | 1042 | -5905794062536720418
Waikanae       | 5          | 1025 | -5905794062536720418
Hutt Valley Line | 11         | 2025 | -2183064056535108044
Hutt Valley Line | 11         | 2019 | -2183064056535108044
Hutt Valley Line | 11         | 2010 | -2183064056535108044
Hutt Valley Line | 11         | 2001 | -2183064056535108044
Hutt Valley Line | 11         | 1955 | -2183064056535108044
Hutt Valley Line | 11         | 1952 | -2183064056535108044
Hutt Valley Line | 11         | 1947 | -2183064056535108044
Hutt Valley Line | 11         | 1935 | -2183064056535108044
Hutt Valley Line | 2          | 1045 | 2322329569350831795
Hutt Valley Line | 2          | 1033 | 2322329569350831795
Hutt Valley Line | 2          | 1025 | 2322329569350831795
Hutt Valley Line | 2          | 1015 | 2322329569350831795
Hutt Valley Line | 2          | 1000 | 2322329569350831795

(30 rows)
```

With filter for our specific data.

```
$ ccm node1 cqlsh -e "USE ass2; SELECT line_name, service_no, time, token(line_name, service_no) as t FROM time_table WHERE line_name = 'Hutt Valley Line';"

line_name      | service_no | time | t
-----+-----+-----+-----
Hutt Valley Line | 2          | 1045 | 2322329569350831795

(1 rows)
```

Because we still have our node ring details, we can find the first node where this token belongs.

dc1:						
127.0.0.4	r1	Up	Normal	98.97 KiB	20.00%	1844674407370955161
127.0.0.5	r1	Up	Normal	98.96 KiB	20.00%	5534023222112865484
dc2:						
127.0.0.8	r1	Up	Normal	98.94 KiB	25.00%	100
127.0.0.9	r1	Up	Normal	98.97 KiB	25.00%	4611686018427388004

As we can see in node ring list our "Hutt Valley Line" token is smaller then the threshold of `node5` of `dc1` . So our record default node is `node5` . Replicas are the following in the ring: `node1` , `node2`   
In `dc2` token number is smaller than the threshold of `node9` , the primary node will be `node9` and replicas will be stored on `node6` and `node7` , because they are the following nodes in the ring of `dc2` .

Of course, we have to test our solution, so when we get back our tools, we can test above numbers with a "brute force" way, as we did in Question 6. Stop nodes, use `consistency one` to determine which node responds and which not.

```

$ ccm stop
$ ccm node9 start
$ ccm status
Cluster: 'multi_dc'
-----
node9: UP
node8: DOWN
node1: DOWN
node3: DOWN
node2: DOWN
node5: DOWN
node4: DOWN
node7: DOWN
node6: DOWN
$ ccm node9 cqlsh -e "CONSISTENCY LOCAL_ONE; SELECT * FROM ass2.time_table WHERE line_name='Hutt Valley Line' AND
Consistency level set to LOCAL_ONE.

 line_name      | service_no | time | distance | latitude | longitude | stop
-----+-----+-----+-----+-----+-----+-----
Hutt Valley Line |          2 | 1045 |    34.3 | -41.2865 |  174.7762 | Wellington

(1 rows)

```

It shows, that `node9` stores our requested record.

Repeating the above steps for other nodes, I got respond from `node6` and `node7`, but not from `node8`.

And testing `dc1`, I got respond from `node1`, `node2` and `node5`, but not from `node3` and `node4`.

So our token solution was right. (Automated bash script is saved in `q15-brute-force.sh`)